





How uniform is the average untrained model?

30 random initializations per architecture · 50,257 vocabulary tokens



Each token is averaged across seeds FIRST; plots then show the spread across vocabulary tokens.
1 × uniform = 0.00199% probability. Bars below are token percentiles, not confidence intervals.
Remaining spread includes finite-seed noise; it does not by itself establish a systematic token bias.